

# To do list

- Conda environment
- notebooks on github
- Gitignore datasets
- READ ME
- Instructions for getting data

Andrew Long • [awlong20@gmail.com](mailto:awlong20@gmail.com) • [linkedin.com/in/awlong/](https://www.linkedin.com/in/awlong/)  
[https://github.com/andrewlong/odsc\\_west\\_2018](https://github.com/andrewlong/odsc_west_2018)



# OPEN DATA SCIENCE CONFERENCE



@ODSC


San Francisco | October 31 - Nov. 3 2018



# Introduction to Clinical Natural Language Processing

Andrew Long

<https://towardsdatascience.com/introduction-to-clinical-natural-language-processing-predicting-hospital-readmission-with-1736d52bc709>



Predicting Hospital Readmission  
with Discharge Summaries

Andrew Long • [awlong20@gmail.com](mailto:awlong20@gmail.com) • [linkedin.com/in/awlong/](https://www.linkedin.com/in/awlong/)  
[https://github.com/andrewlong/odsc\\_west\\_2018](https://github.com/andrewlong/odsc_west_2018)





## FRESENIUS MEDICAL CARE



180,000+  
U.S. PATIENTS  
SERVED



26M  
ANNUAL  
HEMODIALYSIS  
TREATMENT




50+  
STATES AND  
TERRITORIES IN OUR  
NETWORK




2,200+  
U.S. DIALYSIS CLINICS



60,000+  
U.S. EMPLOYEES





# Clinical Notes

- Chest pain
- Shortness of breathe
- Nausea, vomiting, diarrhea
- Weakness
- Sick
- 
- 
- 

Build predictive models that incorporate free-text clinical notes

# Workshop Overview

- Brief overview of clinical dataset (MIMIC III)
  - How to prepare data for a machine learning project
  - How to preprocess the unstructured notes
  - How to build a simple predictive model using a bag-of-words approach
  - How to assess the quality of your model
  - How to decide the next step for improving the model
- 
- Note: I created an artificial dataset based on Stanford's IMDB which you can use if you don't have MIMIC access for the workshop

# Workshop Project Question

## Scalable and accurate deep learning for electronic health records

Rajkomar et al. (paper at <https://arxiv.org/abs/1801.07860>)


- in-hospital mortality (AUC = 0.93–0.94)
- 30-day unplanned readmission (AUC = 0.75–76)
- prolonged length of stay (AUC = 0.85–0.86)
- discharge diagnoses (AUC = 0.90)

AUC is a data science performance metric where closer to 1 is better

How good of a model can we get if use the discharge free-text summaries with a simple predictive model to predict readmission?

# Classification Model Definition

- Predict which patients are at risk for 30-day unplanned readmission utilizing free-text hospital discharge summaries.





# Part 0: MIMIC III dataset

Andrew Long • [awlong20@gmail.com](mailto:awlong20@gmail.com) • [linkedin.com/in/awlong/](https://www.linkedin.com/in/awlong/)  
[https://github.com/andrewwlong/odsc\\_west\\_2018](https://github.com/andrewwlong/odsc_west_2018)



# MIMIC III dataset

- This database contains de-identified data from over 40,000 patients who were admitted to Beth Israel Deaconess Medical Center in Boston, Massachusetts from 2001 to 2012

Access:

<https://mimic.physionet.org/gettingstarted/access/>

<https://towardsdatascience.com/getting-access-to-mimic-iii-hospital-database-for-data-science-projects-791813feb735>

- Since dataset has restricted access, any single subject data shown in this workshop is artificially created.

Andrew Long • [awlong20@gmail.com](mailto:awlong20@gmail.com) • [linkedin.com/in/awlong/](https://www.linkedin.com/in/awlong/)  
[https://github.com/andrewlong/odsc\\_west\\_2018](https://github.com/andrewlong/odsc_west_2018)



# MIMIC III datasets for tutorial

<https://physionet.org/works/MIMICIIIClinicalDatabase>

- ADMISSIONS.csv.gz
- NOTEEVENTS.csv.gz
- placed in a 'data' folder in same folder as this workshop's notebook

```
import gzip

for filename in ["data/ADMISSIONS.csv.gz", "data/NOTEVENTS.csv.gz"]:

    with gzip.open(filename, 'rt') as f:
        data = f.read()
    with open(filename[:-3], 'wt') as f:
        f.write(data)
```

1. [checksum\\_md5\\_zipped.txt](#) (MD5 checksum for zipped files)
2. [checksum\\_md5\\_unzipped.txt](#) (MD5 checksum for unzipped files)
3. [ADMISSIONS.csv.gz](#) (2.5M compressed, 12M decompressed)
4. [CALLOUT.csv.gz](#) (1.2M compressed, 6.1M decompressed)
5. [CAREGIVERS.csv.gz](#) (49K compressed, 199K decompressed)
6. [CHARTEVENTS.csv.gz](#) (4.0G compressed, 33G decompressed)
7. [CPTEVENTS.csv.gz](#) (4.8M compressed, 56M decompressed)
8. [DATETIMEEVENTS.csv.gz](#) (53M compressed, 502M decompressed)
9. [DIAGNOSES\\_ICD.csv.gz](#) (4.5M compressed, 19M decompressed)
10. [DRGCODES.csv.gz](#) (1.7M compressed, 11M decompressed)
11. [D\\_CPT.csv.gz](#) (3.9K compressed, 14K decompressed)
12. [D\\_ICD\\_DIAGNOSES.csv.gz](#) (279K compressed, 1.4M decompressed)
13. [D\\_ICD\\_PROCEDURES.csv.gz](#) (75K compressed, 305K decompressed)
14. [D\\_ITEMS.csv.gz](#) (184K compressed, 933K decompressed)
15. [D\\_LABITEMS.csv.gz](#) (12K compressed, 43K decompressed)
16. [ICUSTAYS.csv.gz](#) (1.9M compressed, 6.1M decompressed)
17. [INPUTEVENTS\\_CV.csv.gz](#) (403M compressed, 2.3G decompressed)
18. [INPUTEVENTS\\_MV.csv.gz](#) (144M compressed, 931M decompressed)
19. [LABEVENTS.csv.gz](#) (321M compressed, 1.8G decompressed)
20. [MICROBIOLOGYEVENTS.csv.gz](#) (7.3M compressed, 70M decompressed)
21. [NOTEVENTS.csv.gz](#) (1.1G compressed, 3.8G decompressed)
22. [OUTPUTEVENTS.csv.gz](#) (56M compressed, 379M decompressed)
23. [PATIENTS.csv.gz](#) (559K compressed, 2.6M decompressed)
24. [PRESCRIPTIONS.csv.gz](#) (99M compressed, 735M decompressed)
25. [PROCEDUREEVENTS\\_MV.csv.gz](#) (7.5M compressed, 47M decompressed)
26. [PROCEDURES\\_ICD.csv.gz](#) (1.8M compressed, 6.5M decompressed)
27. [SERVICES.csv.gz](#) (1.2M compressed, 3.4M decompressed)
28. [TRANSFERS.csv.gz](#) (5.3M compressed, 24M decompressed)

Andrew Long • [awlong20@gmail.com](mailto:awlong20@gmail.com) • [linkedin.com/in/awlong/](https://www.linkedin.com/in/awlong/)  
[https://github.com/andrewwlong/odsc\\_west\\_2018](https://github.com/andrewwlong/odsc_west_2018)



# Part 1: How to prepare data for a machine learning project

Andrew Long • [awlong20@gmail.com](mailto:awlong20@gmail.com) • [linkedin.com/in/awlong/](https://www.linkedin.com/in/awlong/)  
[https://github.com/andrewlong/odsc\\_west\\_2018](https://github.com/andrewlong/odsc_west_2018)



# Load, clean, merge dataset

```
In [2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

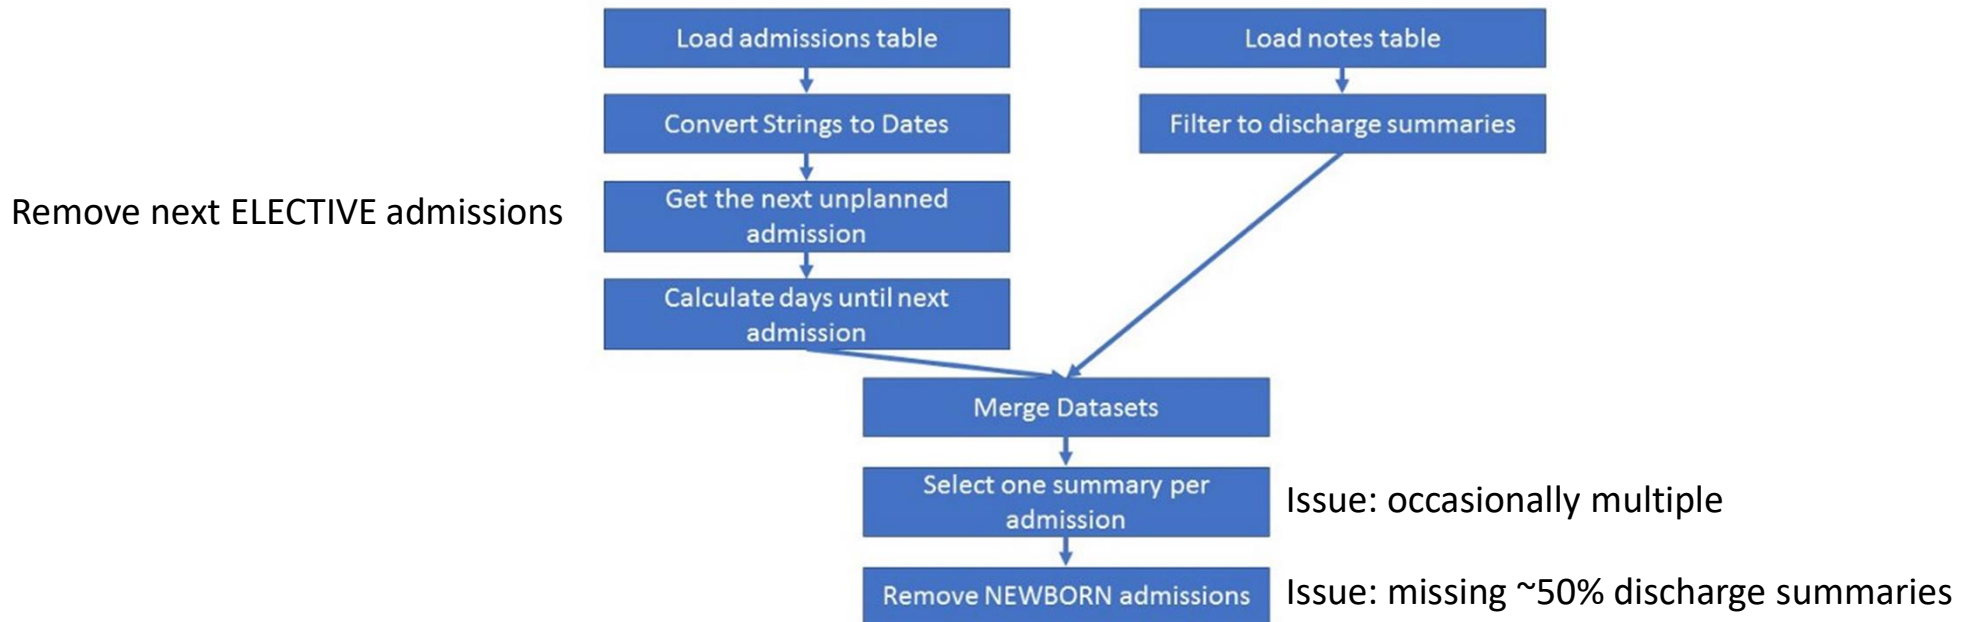
import odsc2018_utils
|
df_adm_notes_clean = odsc2018_utils.load_clean_merge_dataset('data/ADMISSIONS.csv', 'data/NOTEEVENTS.csv')

C:\Users\3236283\AppData\Local\Continuum\Anaconda3\lib\site-packages\IPython\core\interactiveshell.py:2910:
ing: Columns (4,5) have mixed types. Specify dtype option on import or set low_memory=False.
exec(code_obj, self.user_global_ns, self.user_ns)
```

We skip this process to save some time In workshop.

See additional Jupyter Notebook (odsc\_2018\_mimic\_pre) for tutorial on these steps.


# Prepare data for ML Project



# df\_adm\_notes\_clean

- SUBJECT\_ID – unique patient identifier
- HADM\_ID- unique admission identifier
- ADMITTIME – admission date
- DISCHTIME – discharge date
- DEATHTIME – death date
- DAYS\_NEXT\_ADMIT – days until next admission if it exists
- TEXT – discharge summary for this admission

# Add OUTPUT\_LABEL



```
In [3]: df_adm_notes_clean['OUTPUT_LABEL'] = (df_adm_notes_clean.DAYS_NEXT_ADMIT < 30).astype('int')
```

```
In [4]: print('Number of positive samples:', (df_adm_notes_clean.OUTPUT_LABEL == 1).sum())
print('Number of negative samples:', (df_adm_notes_clean.OUTPUT_LABEL == 0).sum())
print('Total samples:', len(df_adm_notes_clean))
```

```
Number of positive samples: 3004
Number of negative samples: 48109
Total samples: 51113
```



# Make training/validation/test sets



- Training samples: these samples are used to train the model
- Validation samples: these samples are held out from the training data and are used to make decisions on how to improve the model
- Test samples: these samples are held out from all decisions and are used to measure the generalized performance of the model

# Make training/validation/test sets

70% Training

15% Validation

15% Test

70/15/15 is a design choice

```
In [5]: # shuffle the samples
df_adm_notes_clean = df_adm_notes_clean.sample(n = len(df_adm_notes_clean), random_state = 42)
df_adm_notes_clean = df_adm_notes_clean.reset_index(drop = True)

# Save 30% of the data as validation and test data
df_valid_test=df_adm_notes_clean.sample(frac=0.30,random_state=42)

df_test = df_valid_test.sample(frac = 0.5, random_state = 42)
df_valid = df_valid_test.drop(df_test.index)

# use the rest of the data as training data
df_train_all=df_adm_notes_clean.drop(df_valid_test.index)
```

# Make training/validation/test sets

70% Training

15% Validation

15% Test

70/15/15 is a design choice

Verify that positive prevalence is approximately the same in the 3 groups

```
In [6]: print('Test prevalence(n = %d):%.3f'%(len(df_test),df_test.OUTPUT_LABEL.sum()/ len(df_test)))
print('Valid prevalence(n = %d):%.3f'%(len(df_valid),df_valid.OUTPUT_LABEL.sum()/ len(df_valid)))
print('Train all prevalence(n = %d):%.3f'%(len(df_train_all), df_train_all.OUTPUT_LABEL.sum()/ len(df_train_all)))
print('all samples (n = %d)'%len(df_adm_notes_clean))
assert len(df_adm_notes_clean) == (len(df_test)+len(df_valid)+len(df_train_all)), 'math didnt work'
```

```
Test prevalence(n = 7667):0.062
Valid prevalence(n = 7667):0.057
Train all prevalence(n = 35779):0.058
all samples (n = 51113)
```


```
In [7]: df_train_all.to_csv('data/df_train_all.csv',index=False)
df_valid.to_csv('data/df_valid.csv',index=False)
df_test.to_csv('data/df_test.csv',index=False)
```

# Artificial Dataset

- For those without MIMIC data, I created an artificial dataset based on IMDB dataset.
  - data\_artificial/df\_train\_all\_imdb.csv
  - data\_artificial/df\_valid\_imdb.csv
  - data\_artificial/df\_test\_imdb.csv
- 
- validation and test sets were created to have approximately same number and prevalence as the MIMIC sets

# Imbalanced Classification

- Model that always guesses 'Not readmitted' → 94% accuracy, but never catches any readmissions (0% recall)
- To prevent this from happening, we need to balance the training set
  - **sub-sample the more dominant class: use a random subset of the negatives**
  - over-sample the imbalanced class: use the same positive samples multiple times
  - create synthetic positive data



# Subsample Training Dataset

```
In [9]: # split the training data into positive and negative
rows_pos = df_train_all.OUTPUT_LABEL == 1
df_train_pos = df_train_all.loc[rows_pos]
df_train_neg = df_train_all.loc[~rows_pos]

n = np.min([len(df_train_pos), len(df_train_neg)])

# merge the balanced data
df_train = pd.concat([df_train_pos.sample(n = n, random_state = 42), \
                      df_train_neg.sample(n = n, random_state = 42)], axis = 0)

# shuffle the order of training samples
df_train = df_train.sample(n = len(df_train), random_state = 42).reset_index(drop = True)

print('Train prevalence (n = %d):' % len(df_train), df_train.OUTPUT_LABEL.sum() / len(df_train))

Train prevalence (n = 4184): 0.5
```

70% Training

Positive

Negative

# Part 2: How to preprocess the unstructured notes

Andrew Long • [awlong20@gmail.com](mailto:awlong20@gmail.com) • [linkedin.com/in/awlong/](https://www.linkedin.com/in/awlong/)  
[https://github.com/andrewlong/odsc\\_west\\_2018](https://github.com/andrewlong/odsc_west_2018)



# Pre-process Text Data

- Occasionally, need to modify the text to make useable (for example drop newlines, carriage returns, numbers, etc)
- Two Methods:
  - Modify the original dataframe TEXT column
  - Pre-process as part of the pipeline



# Modify Original Text

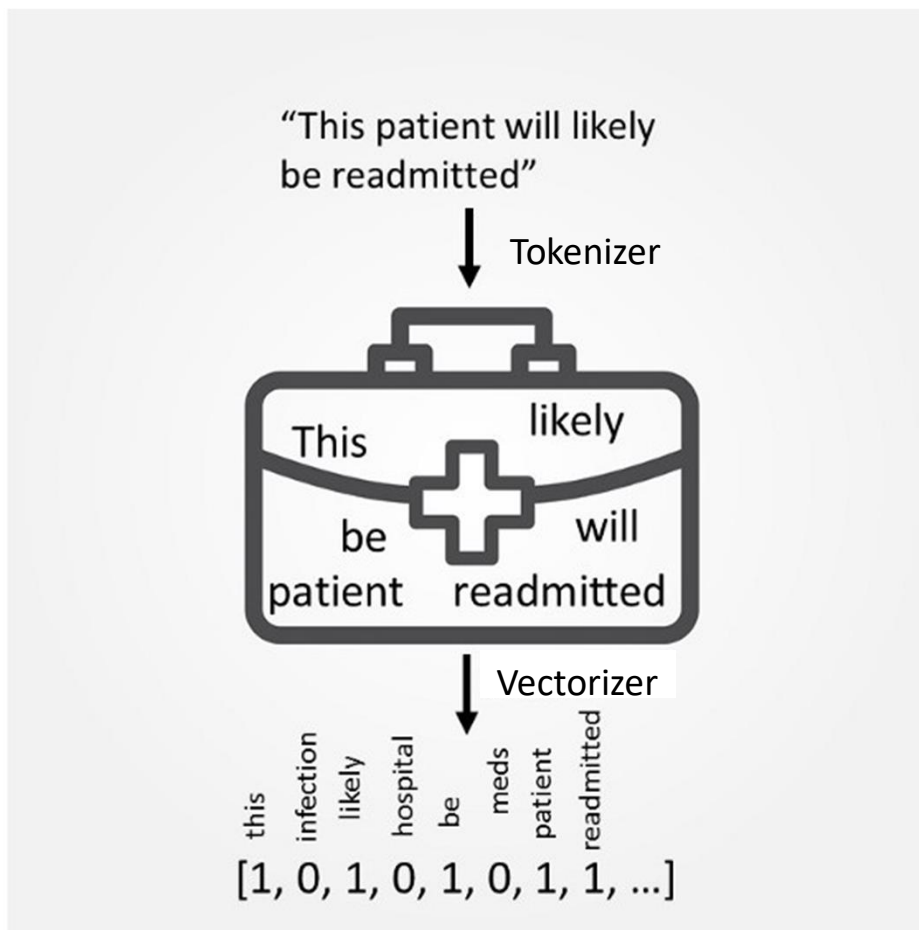
Drop newline, carriage returns  
Replace missing notes with a space

```
In [12]: def preprocess_text(df):  
        # This function preprocesses the text by filling  
        df.TEXT = df.TEXT.fillna(' ')  
        df.TEXT = df.TEXT.str.replace('\n', ' ')  
        df.TEXT = df.TEXT.str.replace('\r', ' ')  
        return df
```

```
In [13]: # preprocess the text to deal with known issues  
df_train = preprocess_text(df_train)  
df_valid = preprocess_text(df_valid)  
df_test = preprocess_text(df_test)
```

# Bag-of-words

- Split a note into tokens (i.e. words) then 'count' the number of each token
- Use these 'counts' as feature columns
- Note: different techniques for 'counts'



# Bag-of-words design choices

- How to preprocess the words into tokens
- How to count the tokens
- Which tokens to use

# Build a tokenizer

```
In [14]: import nltk
         from nltk import word_tokenize
         word_tokenize('This should be tokenized. 11/01/2018 sentence has stars**')

Out[14]: ['This',
          'should',
          'be',
          'tokenized',
          '.',
          '11/01/2018',
          'sentence',
          'has',
          'stars**']
```

- Sentence is tokenized by spaces and some punctuation but not all.
- Numbers are also still included
- 'This' would be considered a different token than 'this'
- 'stars\*\*' would be different than 'stars'

# Build a custom tokenizer

- Replace punctuation with spaces
- Replace numbers with spaces
- Lowercase all words

```
In [15]: import string
print(string.punctuation)

!"#$%&'()*+,-./:;<=>?@[\\]^_`{|}~
```

```
In [16]: def tokenizer_better(text):
# tokenize the text by replacing punctuation and numbers with spaces and lowercase all words

punc_list = string.punctuation+'0123456789'
t = str.maketrans(dict.fromkeys(punc_list, " "))
text = text.lower().translate(t)
tokens = word_tokenize(text)
return tokens
```

Fast way to replace characters with spaces

```
In [17]: tokenizer_better('This should be tokenized. 11/01/2018 sentence has stars**')
```

```
Out[17]: ['this', 'should', 'be', 'tokenized', 'sentence', 'has', 'stars']
```

# Build a simple vectorizer

```
In [18]: sample_text = ['Open Data Science Conference is about learning',  
                        'Data data DATA',  
                        'Learning is part of data science']
```

```
In [19]: from sklearn.feature_extraction.text import CountVectorizer  
vect = CountVectorizer(tokenizer = tokenizer_better)  
vect.fit(sample_text)  
  
# matrix is stored as a sparse matrix (since you have a lot of zeros)  
X = vect.transform(sample_text)
```

Specify custom tokenizer



- CountVectorizer is the simplest method for bag-of-words
- Counts the number of occurrences of each word
- Other common method is TfidfVectorizer which takes into account frequency of word usage across notes

# Build a simple vectorizer

```
In [18]: sample_text = ['Open Data Science Conference is about learning',  
                        'Data data DATA',  
                        'Learning is part of data science']
```

```
In [19]: from sklearn.feature_extraction.text import CountVectorizer  
vect = CountVectorizer(tokenizer = tokenizer_better)  
vect.fit(sample_text)  
  
# matrix is stored as a sparse matrix (since you have a lot of zeros)  
X = vect.transform(sample_text)
```

```
In [20]: X
```

```
Out[20]: <3x9 sparse matrix of type '<class 'numpy.int64'>'  
         with 14 stored elements in Compressed Sparse Row format>
```

```
In [21]: # we can visualize this small example if we convert it to an array  
X.toarray()
```

```
Out[21]: array([[1, 1, 1, 1, 1, 0, 1, 0, 1],  
                [0, 0, 3, 0, 0, 0, 0, 0, 0],  
                [0, 0, 1, 1, 1, 1, 0, 1, 1]], dtype=int64)
```

```
: # get the column names  
vect.get_feature_names()
```

```
: ['about',  
   'conference',  
   'data',  
   'is',  
   'learning',  
   'of',  
   'open',  
   'part',  
   'science']
```

# Train clinical vectorizer

```
In [23]: from sklearn.feature_extraction.text import CountVectorizer
vect = CountVectorizer(max_features = 3000, tokenizer = tokenizer_better)

# this could take a while
vect.fit(df_train.TEXT.values)
```

```
Out[23]: CountVectorizer(analyzer='word', binary=False, decode_error='strict',
dtype=<class 'numpy.int64'>, encoding='utf-8', input='content',
lowercase=True, max_df=1.0, max_features=3000, min_df=1,
ngram_range=(1, 1), preprocessor=None, stop_words=None,
strip_accents=None, token_pattern='(?u)\\b\\w\\w+\\b',
tokenizer=<function tokenizer_better at 0x000001CA1DA7A400>,
vocabulary=None)
```

Good practice to specify the max\_features (otherwise it could take a long time with big data set)  
Size of max\_features is then a hyperparameter to tune



# Stop words

“the”, “is”, “are”, “and”

- Stop word – commonly used words with little value to ML model
- Frequency of word use depends on domain (clinical, twitter, Wikipedia)

# Stop words

<https://www.linkedin.com/pulse/another-twitter-sentiment-analysis-python-part-2-ricky-kim/>

```
In [24]: neg_doc_matrix = vect.transform(df_train[df_train.OUTPUT_LABEL == 0].TEXT)
pos_doc_matrix = vect.transform(df_train[df_train.OUTPUT_LABEL == 1].TEXT)
neg_tf = np.sum(neg_doc_matrix,axis=0)
pos_tf = np.sum(pos_doc_matrix,axis=0)
neg = np.squeeze(np.asarray(neg_tf))
pos = np.squeeze(np.asarray(pos_tf))

term_freq_df = pd.DataFrame([neg,pos],columns=vect.get_feature_names()).transpose()
term_freq_df.columns = ['negative', 'positive']
term_freq_df['total'] = term_freq_df['negative'] + term_freq_df['positive']
term_freq_df.sort_values(by='total', ascending=False).iloc[:10]
```

Out[24]:

	negative	positive	total
the	71054	76756	147810
and	62455	71658	134113
to	53226	62085	115311
of	51303	60491	111794
was	48074	53521	101595
with	38036	44583	82619
a	35428	41629	77057
on	32290	39765	72055
mg	27718	39045	66763
in	29567	34755	64322


More medications predictive of readmission?

Fast technique for  
finding term frequency

Andrew Long • [awlong20@gmail.com](mailto:awlong20@gmail.com) • [linkedin.com/in/awlong/](https://www.linkedin.com/in/awlong/)  
[https://github.com/andrewwlong/odsc\\_west\\_2018](https://github.com/andrewwlong/odsc_west_2018)



# Stop words



```
In [26]: my_stop_words = ['the', 'and', 'to', 'of', 'was', 'with', 'a', 'on', 'in', 'for', 'name',  
                          'is', 'patient', 's', 'he', 'at', 'as', 'or', 'one', 'she', 'his', 'her', 'am',  
                          'were', 'you', 'pt', 'pm', 'by', 'be', 'had', 'your', 'this', 'date',  
                          'from', 'there', 'an', 'that', 'p', 'are', 'have', 'has', 'h', 'but', 'o',  
                          'namepattern', 'which', 'every', 'also']
```

# Build a vectorizer removing stop words

```
In [27]: from sklearn.feature_extraction.text import CountVectorizer
vect = CountVectorizer(max_features = 3000,
                      tokenizer=tokenizer_better,
                      stop_words = my_stop_words)
# this could take a while
vect.fit(df_train.TEXT.values)
```

```
Out[27]: CountVectorizer(analyzer='word', binary=False, decode_error='strict',
                        dtype=<class 'numpy.int64'>, encoding='utf-8', input='content',
                        lowercase=True, max_df=1.0, max_features=3000, min_df=1,
                        ngram_range=(1, 1), preprocessor=None,
                        stop_words=['the', 'and', 'to', 'of', 'was', 'with', 'a', 'on', 'in', 'for', 'name', 'is', 'patient', 's', 'h
e', 'at', 'as', 'or', 'one', 'she', 'his', 'her', 'am', 'were', 'you', 'pt', 'pm', 'by', 'be', 'had', 'your', 'this',
'date', 'from', 'there', 'an', 'that', 'p', 'are', 'have', 'has', 'h', 'but', 'o', 'namepattern', 'which', 'every', '
also'],
                        strip_accents=None, token_pattern='(?u)\\b\\w\\w+\\b',
                        tokenizer=<function tokenizer_better at 0x000001CA1DA7A400>,
                        vocabulary=None)
```

## Create X, y

```
In [28]: X_train_tf = vect.transform(df_train.TEXT.values)
X_valid_tf = vect.transform(df_valid.TEXT.values)
```

Get labels

```
In [29]: y_train = df_train.OUTPUT_LABEL
y_valid = df_valid.OUTPUT_LABEL
```

# Part 3: How to build a simple predictive model using a bag-of-words approach

Andrew Long • [awlong20@gmail.com](mailto:awlong20@gmail.com) • [linkedin.com/in/awlong/](https://www.linkedin.com/in/awlong/)  
[https://github.com/andrewlong/odsc\\_west\\_2018](https://github.com/andrewlong/odsc_west_2018)



# Logistic Regression

- Traditional Machine Learning algorithm
- Works well with sparse matrices
- Fast to train
- Interpretable

# Logistic Regression

```
In [30]: # logistic regression
from sklearn.linear_model import LogisticRegression
clf=LogisticRegression(C = 0.0001, penalty = 'l2', random_state = 42)
clf.fit(X_train_tf, y_train)

C:\Users\Andy\AppData\Local\conda\conda\envs\odsc_west_2018\lib\site-packag
utureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a
FutureWarning)

Out[30]: LogisticRegression(C=0.0001, class_weight=None, dual=False,
    fit_intercept=True, intercept_scaling=1, max_iter=100,
    multi_class='warn', n_jobs=None, penalty='l2', random_state=42,
    solver='warn', tol=0.0001, verbose=0, warm_start=False)
```

Hyperparameter C is helps control the effect of regularization

We will discuss how to optimize C

Tip: the same C usually doesn't work as well for both CountVectorizer and TfidfVectorizer



# Predictions

```
In [31]: model = clf  
y_train_preds = model.predict_proba(X_train_tf)[: ,1]  
y_valid_preds = model.predict_proba(X_valid_tf)[: ,1]
```

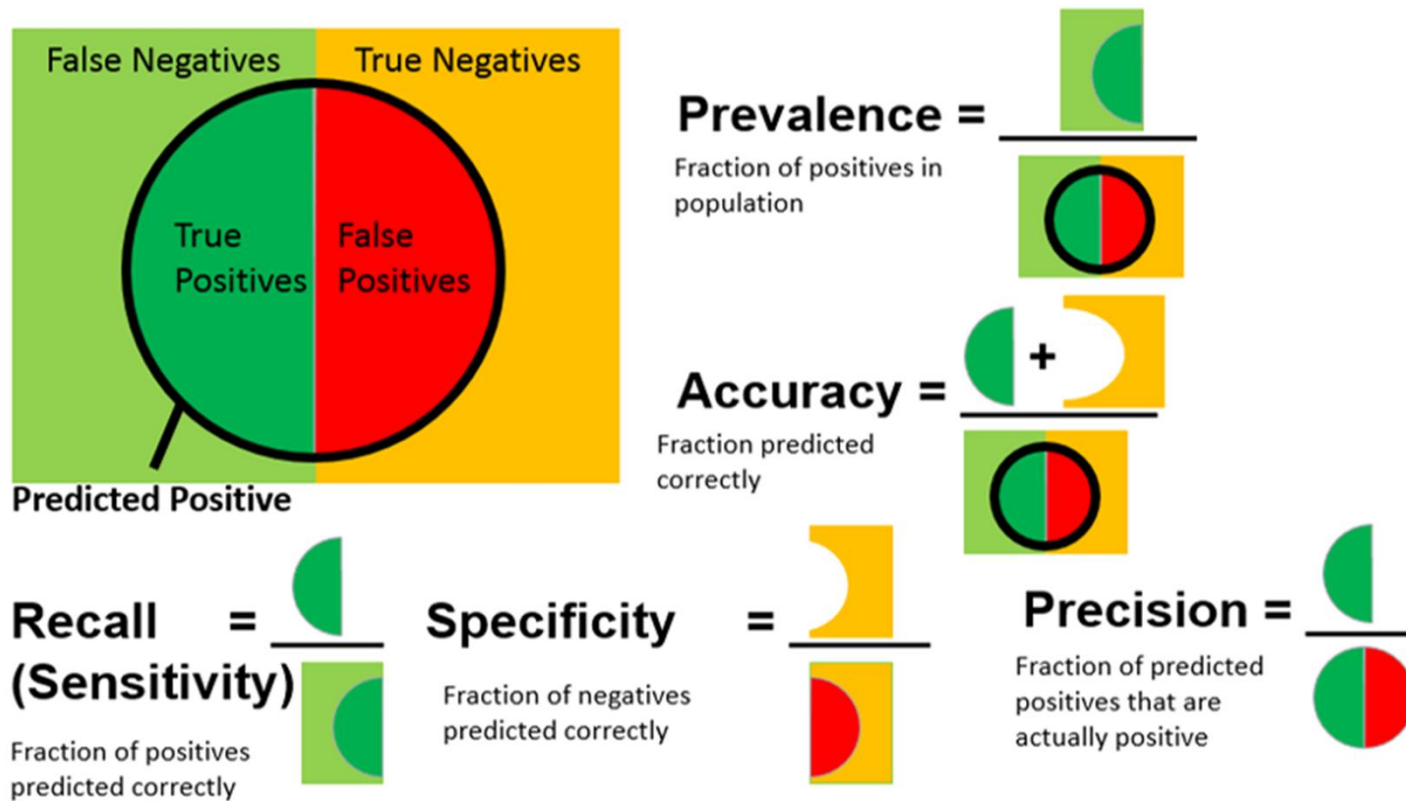
```
In [32]: print(y_train[:10].values)  
print(y_train_preds[:10])  
  
[1 1 0 1 1 1 0 0 1 1]  
[0.76307111 0.63114288 0.29772094 0.77926068 0.59694889 0.55643044  
 0.36292154 0.90984735 0.47806099 0.67622763]
```

# Part 4: How to assess the quality of your model

Andrew Long • [awlong20@gmail.com](mailto:awlong20@gmail.com) • [linkedin.com/in/awlong/](https://www.linkedin.com/in/awlong/)  
[https://github.com/andrewlong/odsc\\_west\\_2018](https://github.com/andrewlong/odsc_west_2018)



# Performance Metrics




**Example: Positive = Hospitalized, Negative = Not Hospitalized**

Andrew Long • [awlong20@gmail.com](mailto:awlong20@gmail.com) • [linkedin.com/in/awlong/](https://www.linkedin.com/in/awlong/)  
[https://github.com/andrewlong/odsc\\_west\\_2018](https://github.com/andrewlong/odsc_west_2018)


#ODSC

# Performance

Metric	Training	Validation
Prevalence	<b>50%</b>	<b>5.7 %</b>
Accuracy	69.5%	68.2 %
Recall	66.6 %	64.8 %
Precision	<b>70.6 %</b>	<b>11.0 %</b>
Specificity	72.3 %	68.4 %
Area Under ROC Curve (AUC)	0.757	0.704



Example: Positive = Hospitalized, Negative = Not Hospitalized



# Part 5: How to decide the next step for improving the model

Andrew Long • [awlong20@gmail.com](mailto:awlong20@gmail.com) • [linkedin.com/in/awlong/](https://www.linkedin.com/in/awlong/)  
[https://github.com/andrewlong/odsc\\_west\\_2018](https://github.com/andrewlong/odsc_west_2018)




# Design decisions


- Which and how much data to use? Should we spend time collecting more data?
- How to tokenize?
  - Should we use stemming? (“stemming” → “stem”)
- How to vectorizer?
  - Change number of words?
  - Switch to tfidfvectorizer?
- How to select hyperparameters in Logistic regression?
- Should we switch to a different ML model?



# Learning Curve (diagnose bias/variance)




High Bias




High Variance


# Learning Curve (diagnose bias/variance)



High Bias



High Variance



Both?

Adding more samples  
Probably won't help majorly  
at this point



# Helpful techniques

- Techniques for reducing bias (underfitting)
  - Add new features
  - Increase model complexity
  - Reduce regularization
  - Change model architecture
- Techniques for reducing variance (overfitting)
  - Add more samples
  - Add regularization
  - Reduce number of features
  - Decrease model complexity
  - Add better features
  - Change model architecture

# Helpful techniques


- Techniques for reducing bias (underfitting)
  - Add new features
  - Increase model complexity
  - Reduce regularization
  - Change model architecture
- Techniques for reducing variance (overfitting)
  - Add more samples
  - Add regularization
  - Reduce number of features
  - Decrease model complexity
  - Add better features
  - Change model architecture

# Feature Importance

Negative Feature Importance Score - Logistic Regression



Positive Feature Importance Score - Logistic Regression



More Negative

Logistic regression coefficient

More positive

Andrew Long • [awlong20@gmail.com](mailto:awlong20@gmail.com) • [linkedin.com/in/awlong/](https://www.linkedin.com/in/awlong/)

[https://github.com/andrewlong/odsc\\_west\\_2018](https://github.com/andrewlong/odsc_west_2018)

#ODSC

# Helpful techniques

- Techniques for reducing bias (underfitting)
  - Add new features
  - Increase model complexity
  - Reduce regularization
  - Change model architecture
- Techniques for reducing variance (overfitting)
  - Add more samples
  - Add regularization
  - Reduce number of features
  - Decrease model complexity
  - Add better features
  - Change model architecture

# Hyper parameter tuning

Regularization C


Stemming

(hospitalizations → hospital)

CountVectorizer or  
Tfidfvectorizer

max\_features


N grams



Source: <https://airandspace.si.edu/multimedia-gallery/5128hjpg>


Andrew Long • [awlong20@gmail.com](mailto:awlong20@gmail.com) • [linkedin.com/in/awlong/](https://www.linkedin.com/in/awlong/)

[https://github.com/andrewlong/odsc\\_west\\_2018](https://github.com/andrewlong/odsc_west_2018)




Source: [https://en.wikipedia.org/wiki/Apollo\\_\(spacecraft\)](https://en.wikipedia.org/wiki/Apollo_(spacecraft))

# Hyper parameter tuning



Higher C = more overfitting



Higher max\_features = more overfitting

# Model Architecture

- Naïve Bayes
- Neural Networks (CNN, RNN)  
with Word2Vec

Andrew Long • [awlong20@gmail.com](mailto:awlong20@gmail.com) • [linkedin.com/in/awlong/](https://www.linkedin.com/in/awlong/)  
[https://github.com/andrewlong/odsc\\_west\\_2018](https://github.com/andrewlong/odsc_west_2018)


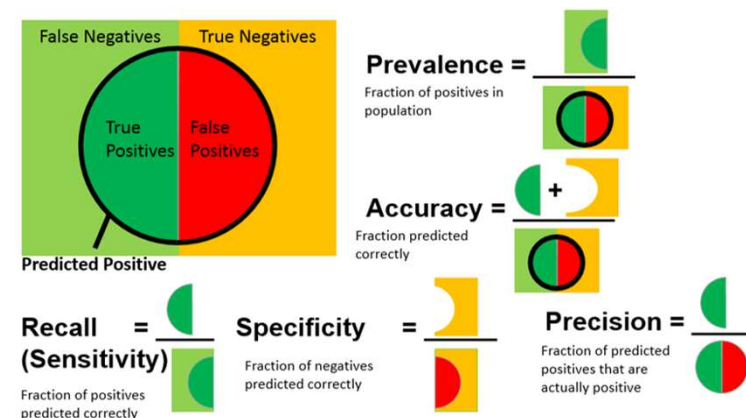


# Final Model (excluding deaths)

Metric	Training	Validation	Test
Prevalence	50%	6.9%	6.6%
Accuracy	0.683	0.672	0.681
Recall	0.644	0.651	0.607
Precision	0.698	0.129	0.120
Specificity	0.722	0.674	0.686
Area Under ROC Curve (AUC)	0.745	0.709	0.704

30-day unplanned readmission (AUC = 0.75–76)  
(Rajkomar et al 2017)

Andrew Long • [awlong20@gmail.com](mailto:awlong20@gmail.com) • [linkedin.com/in/awlong/](https://www.linkedin.com/in/awlong/)  
[https://github.com/andrewwlong/odsc\\_west\\_2018](https://github.com/andrewwlong/odsc_west_2018)



#ODSC






# Introduction to Clinical Natural Language Processing

Andrew Long

<https://towardsdatascience.com/introduction-to-clinical-natural-language-processing-predicting-hospital-readmission-with-1736d52bc709>



Predicting Hospital Readmission  
with Discharge Summaries

Andrew Long • [awlong20@gmail.com](mailto:awlong20@gmail.com) • [linkedin.com/in/awlong/](https://www.linkedin.com/in/awlong/)  
[https://github.com/andrewlong/odsc\\_west\\_2018](https://github.com/andrewlong/odsc_west_2018)

